

FarReach: Write-back Caching in Programmable Switches

Siyuan Sheng¹, Huancheng Puyang¹, Qun Huang², Lu Tang³, and Patrick P. C. Lee¹

¹The Chinese University of Hong Kong ²Peking University ³Xiamen University

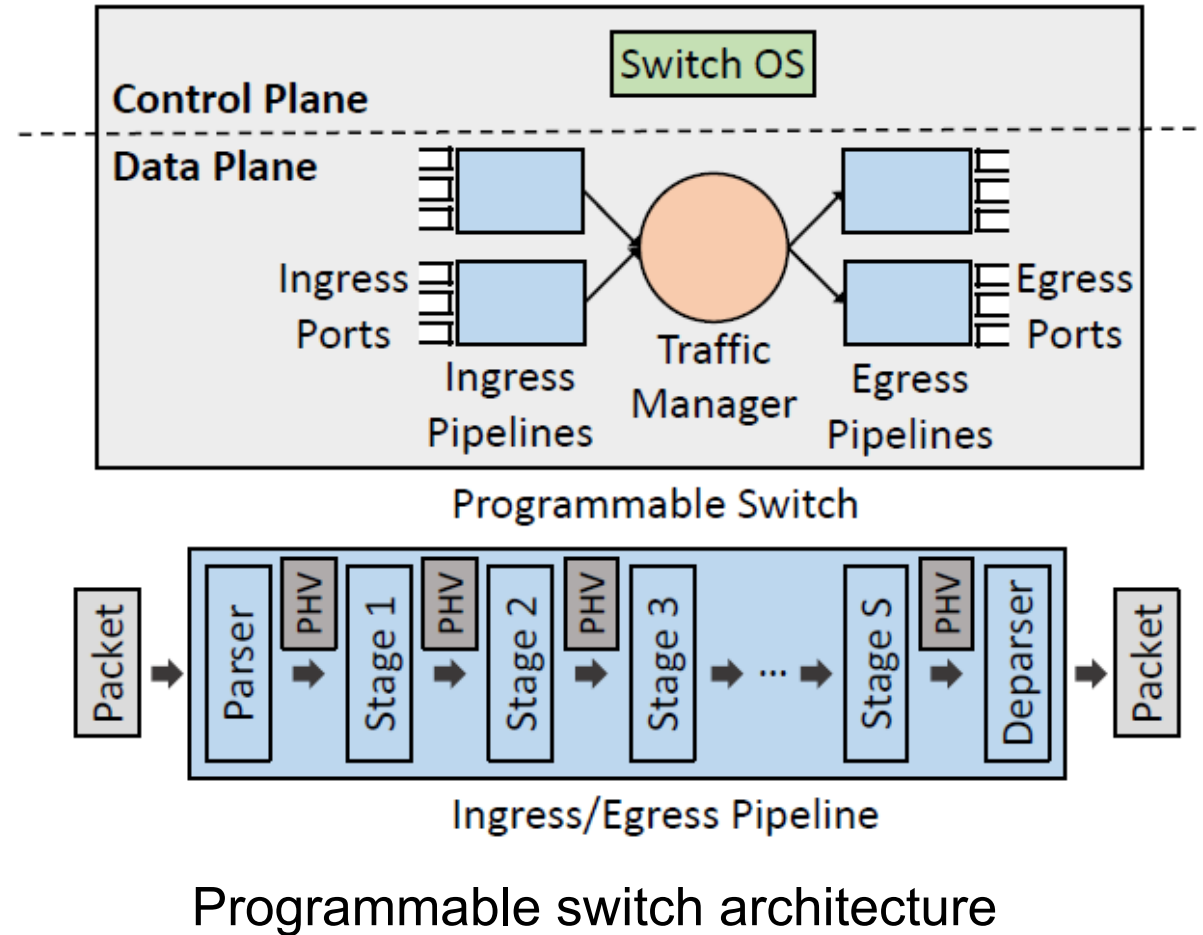
USENIX ATC 2023

Writes in Key-value Stores

- Writes dominate in production key-value storage workloads
 - 20% of Twitter's Twemcache clusters are write-intensive
 - Facebook's RocksDB for AI services has 92.5% of read-modify-writes
- Challenges for high write performance
 - High round-trip latencies in transmission, queuing, and processing
 - Skewness introduces imbalanced server loads

In-switch Write-back Cache

- **Programmable switches** can help improve write performance
 - Switch OS controls multi-pipeline data plane
 - Each pipeline has multiple stages with stateful memory
- **Write-back policy:** caches popular write records in switch without immediately updating servers



Challenges

➤ Performance challenge

- Scarce switch resources require offloading cache management to controller → high controller-to-switch latency

➤ Availability challenge

- Synchronization between switch and servers is required to keep latest records available

➤ Reliability challenge

- Latest records may be lost in switch failures under write-back caching

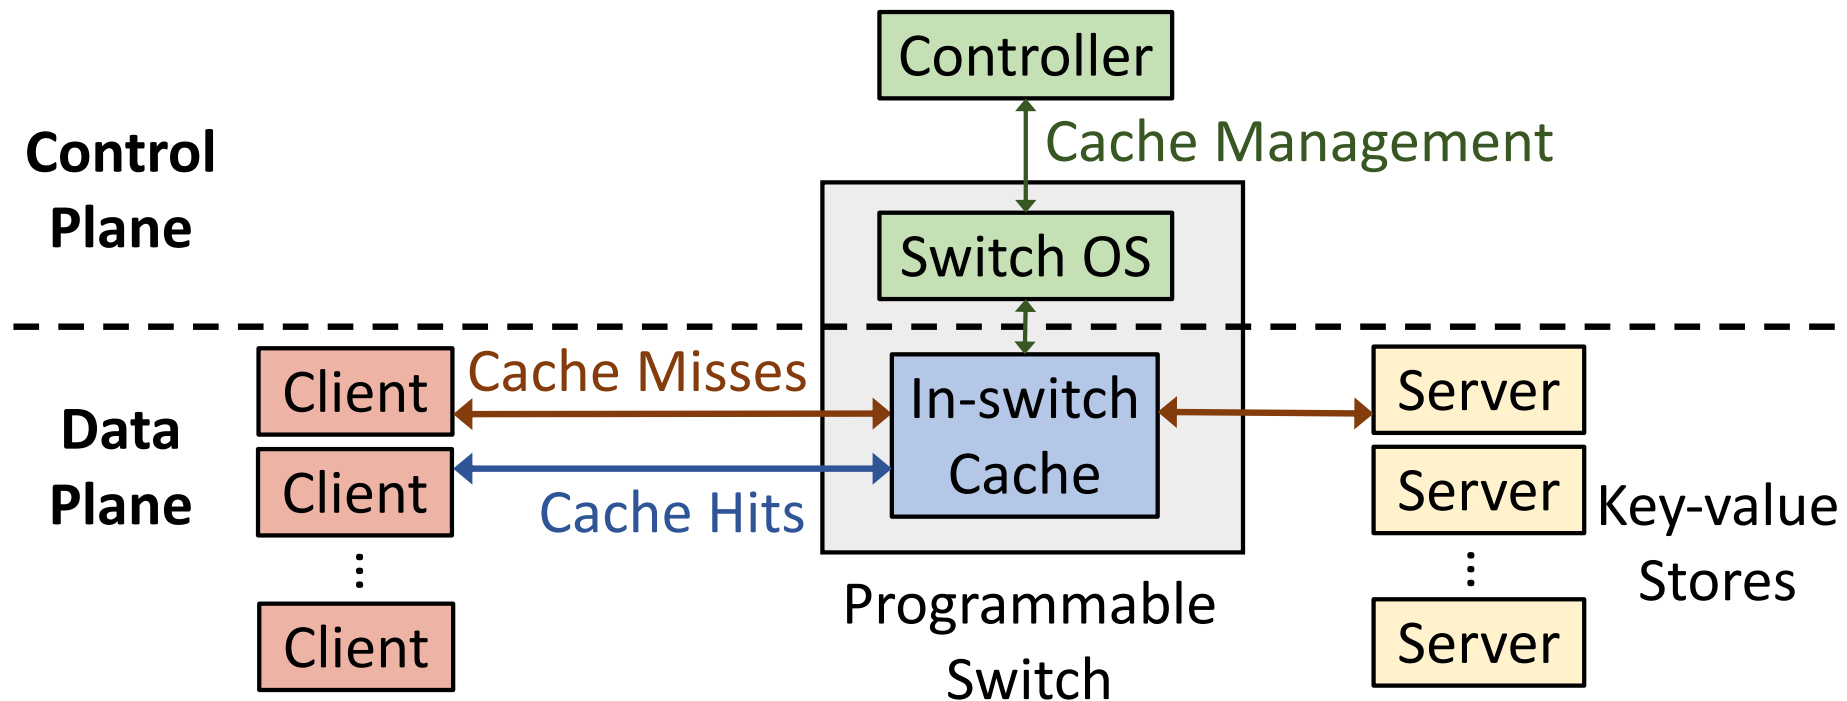
Our Contributions

- **FarReach**, a fast, available, and reliable in-switch write-back cache
 - Non-blocking cache admission for fast access
 - Available cache eviction
 - Crash-consistent snapshot generation and zero-loss recovery
- Prototype implementation
 - P4-based in-switch cache and RocksDB-based servers
- Tofino switch evaluation
 - Up to 6.6× throughput gain under 128 simulated servers
- Open-source FarReach prototype

Design overview

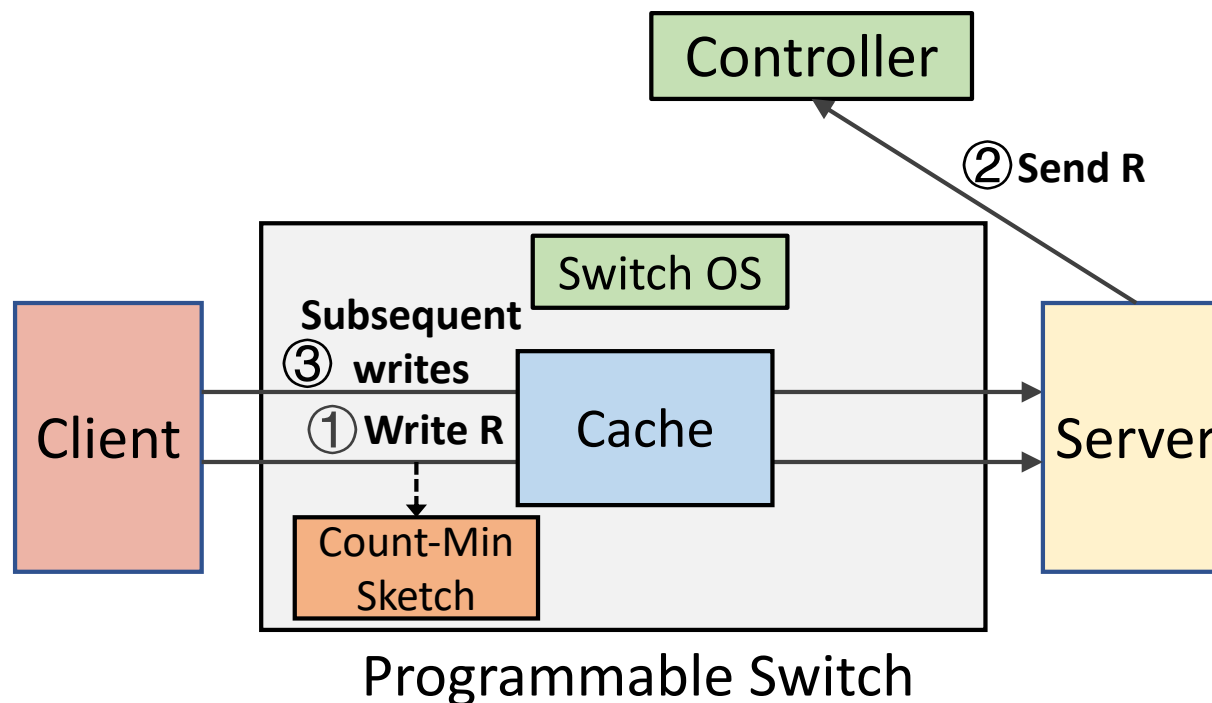
➤ FarReach architecture

- In-switch cache absorbs writes with cache hits
- Controller performs cache management through switch OS
- Carefully co-design control and data planes



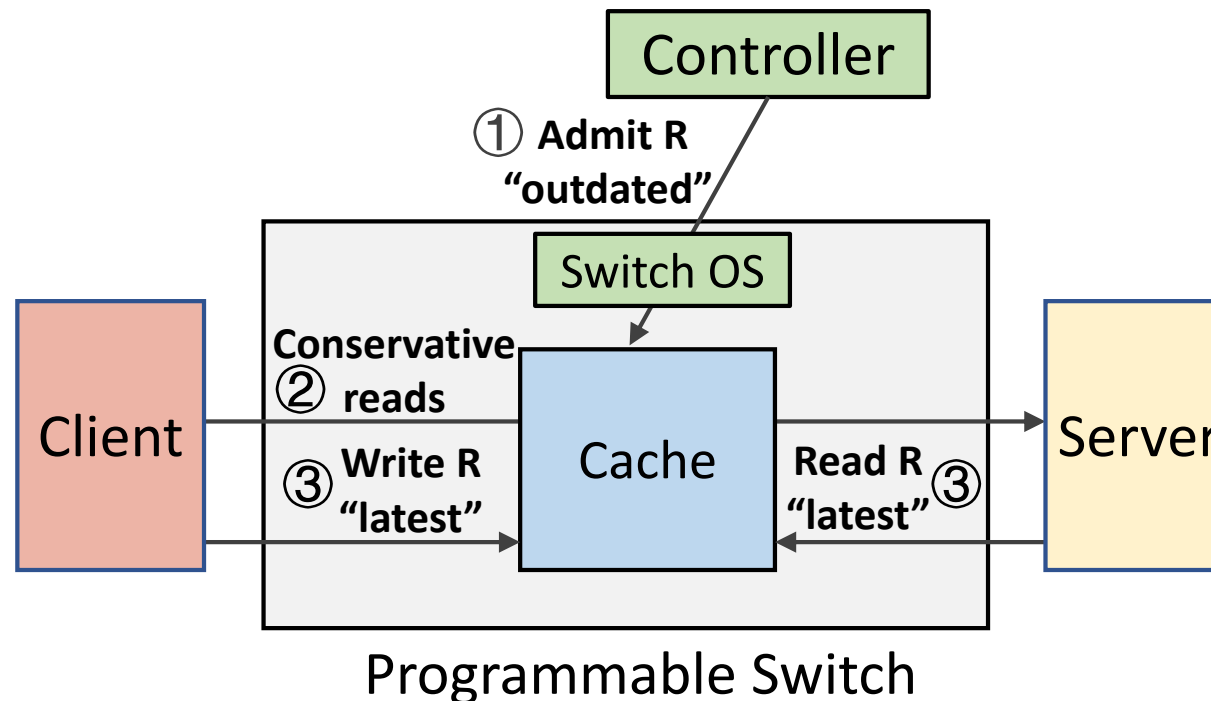
Problem of Cache Admission

- Suppose that a request triggers cache admission
 - Subsequent writes arrive at switch before admission
 - Blocking subsequent writes undermines I/O performance
 - Absorbing subsequent writes in switch undermines availability



Non-blocking Cache Admission

- Process subsequent writes in server without blocking
 - Mark admitted record as “outdated” as server is latest
 - Conservatively forward subsequent reads to server for availability
 - Mark admitted record as “latest” as early as possible




Problem of Cache Eviction

- Under write-back policy
 - Evicted record is latest yet not updated to server
 - Controller loads evicted record to server for persistent storage
- Subsequent writes arrive at switch during cache eviction
 - Processing without synchronization undermines availability
 - Synchronization by controller incurs large overhead

Available Cache Eviction

- Associate additional in-switch metadata to evicted record
 - Mark evicted record as “to-be-evicted”
 - Load evicted record to server before removing it from switch
 - Mark “to-be-evicted” record as “outdated” and forward writes to server
 - Process reads by switch if “latest” or server if “outdated”

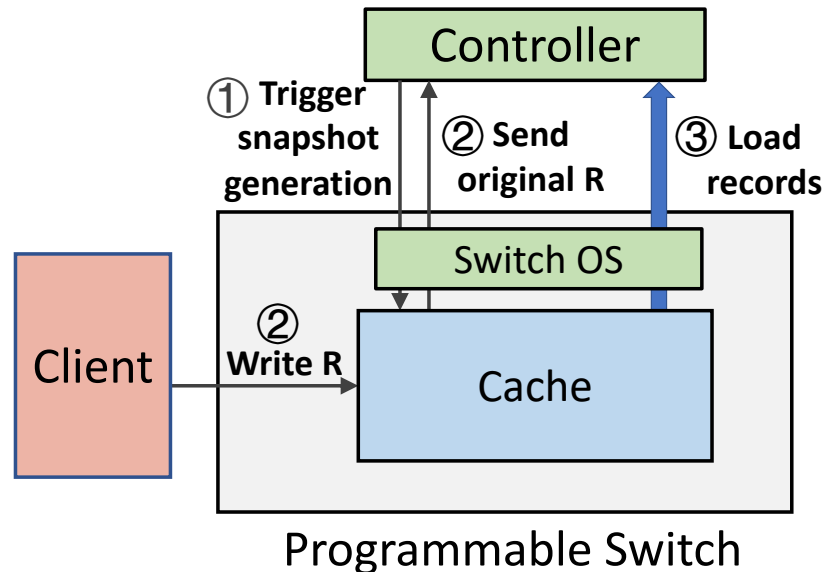


Problem of Reliability

- Under write-back policy
 - Cached records are latest yet not updated to servers
 - Latest in-switch records are lost after switch failures
- Controller loads cached records for snapshots
 - Subsequent writes arrive at switch during snapshot generation
 - Updating cache records incurs inconsistent snapshots

Crash-consistent Snapshot Generation

- Send original cached record for each first write
 - Controller replaces overwritten records for consistency
- Two-phase algorithm
 - Controller triggers snapshot generation
 - Controller loads cached records and switch sends original ones



Zero-loss Recovery

- Limitation of snapshot generation
 - Snapshot generation avoids data loss before the latest snapshot
 - Cached records after the latest snapshot are not protected
- Client-side record preservation
 - Clients preserve copies of cached records after the latest snapshot
 - Controller notifies clients to release the snapshotted records
- Replay-based recovery
 - Replay writes of the latest records to update servers
 - Replay admission decisions to recover in-switch cache

Evaluation

➤ Methodology


- Simulate tens of servers by server rotation for server-side storage
- Compile P4 in a Tofino switch for in-switch cache
- Baselines: NoCache and NetCache [Jin et al., SOSP'17]

➤ Experiments

- YCSB core workloads to evaluate throughput, latency, and scalability
- Synthetic workloads to evaluate impact of different parameters
- Performance of snapshot generation and crash recovery time
- Hardware resource usage


Throughput Analysis

- Simulate 16 servers by server rotation
- Larger throughput especially for workload A
 - In-switch write-back cache reduces server-side load



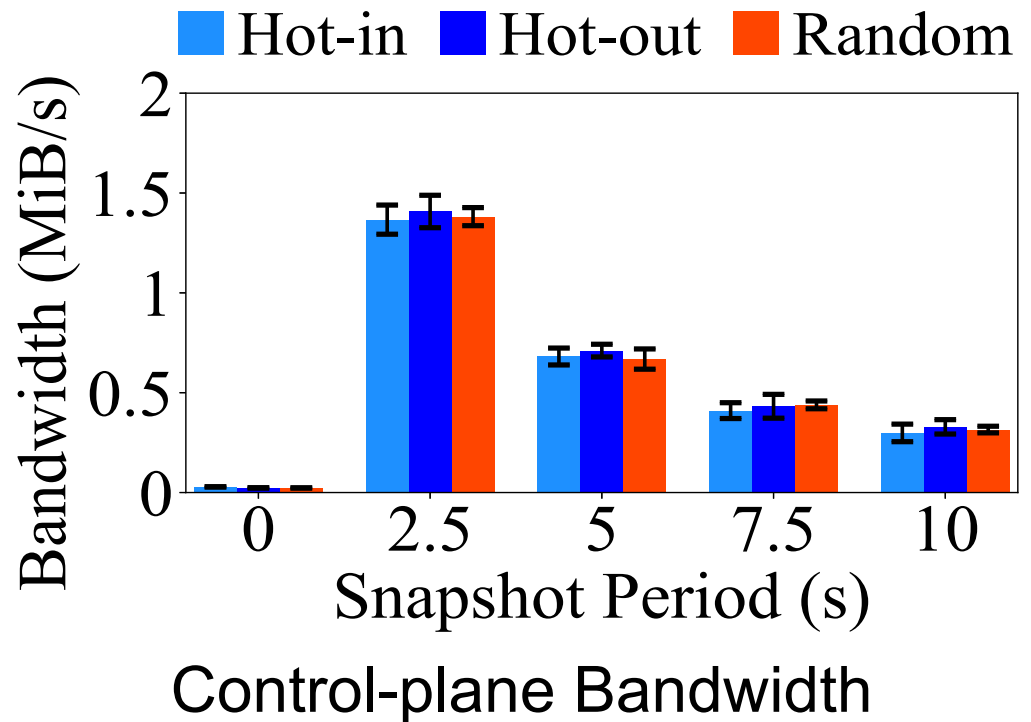
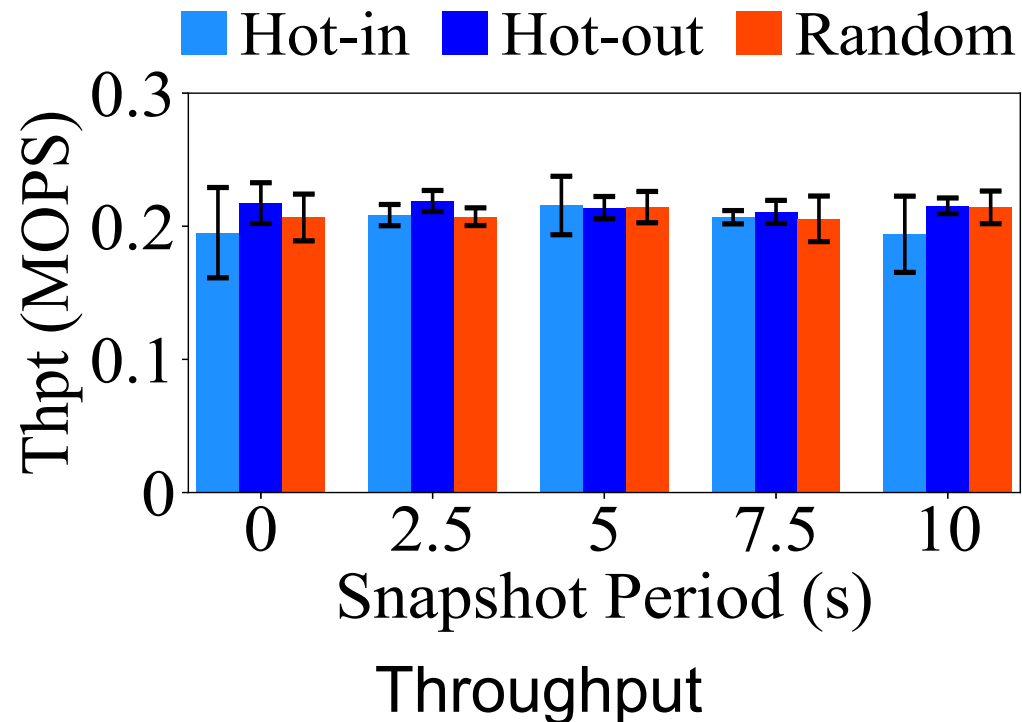
Scalability

- Use workload A (skewed and write-intensive)
 - Simulate 16 to 128 servers by server rotation
- Throughput gain is up to 6.6× under 128 simulated servers
 - In-switch write-back cache balances server-side load



Performance of Snapshot Generation

- Dynamic workload patterns
 - Bandwidth includes snapshot generation and cache management
- Similar throughput and limited control-plane bandwidth



Conclusion

- FarReach, a fast, available, and reliable in-switch write-back cache
 - Non-blocking cache admission
 - Available cache eviction
 - Crash-consistent snapshot generation with zero-loss recovery
- Tofino switch evaluation on YCSB and synthetic workloads
- Source code:
 - <http://adslab.cse.cuhk.edu.hk/software/farreach>

Thank You!
Q & A